



Offensive Language Classification Model

Course: Data Foundations 6713

Team: Automat3rs

Members: Brenda Parnin, Rudy Martinez, Jose Fernandez, Arsalan Bhojani



Agenda

- Background
- Our Approach
- Model Selection Procedure
- Error Analysis



Background

The Connectivity of Social Media

Social Media

- Embedded in the fabric of our society
- Can be leveraged to share aspects of our lives
- Ensures an unfathomable amount of data is available
- Has the potential to unite or destroy





Our Approach

Goal: Create a model that accurately predicts the classification of offensive language in Twitter tweets.

Datasets: Trained and tested models on the train.tsv test.tsv datasets respectively.

Data Identifiers: Twitter identification numbers, Text of Twitter tweets, and Labels that correspond to the tweet's classification.

Data Labels: Not Offensive (NOT), Targeted Insult (TIN), and Untargeted (UNT)

- **NOT:** Posts that do not contain offense or profanity
- **TIN:** Posts containing insult/threat to an individual, a group, or others
- **UNT:** Posts containing non-targeted profanity and swearing. Posts with general profanity are not targeted, but they contain non-acceptable language.



Our Approach Continued...

Process - Models Tested

- LinearSVC, RandomForest, and RandomForest + Lexicon models

Parameters and Features

- A variety of parameters and GridSearchCV values were tested
- A negative word dataset taken from the Data Foundations course was used to create Lexicon Features
- Lexicon Features were incorporated into the Random Forest model to enhance its performance



LinearSVC

Initialize the classifier LinearSVC, Create the params with the C values

```
svc = LinearSVC()
params = {"C": [0.0001, 0.001, 0.01, 0.1, 1., 10., 100.]}
```

Initialize GridSearchCV and Fit the Model

```
clf = GridSearchCV(svc, params, cv = 10, scoring = 'f1_micro')
clf.fit(X_train, y_train)
```

Random Forest

Build Parameter Grid

```
n_estimators = [10, 30, 50, 70, 90, 100]
parameters = {'n_estimators': n_estimators}
```

Initialize GridSearchCV and Fit the Model

```
rand_forest = RandomForestClassifier()
clf_rand = GridSearchCV(rand_forest, parameters, cv = 10, verbose = 2)
clf_rand.fit(X_train, y_train)
```



Lexicon Features

Build Negative Word Lexicon Function

```
class LexiconClassifier():
    def __init__(self):

        self.negative_words = set()
        with open('negative-words.txt', encoding='iso-8859-1') as iFile:
            for row in iFile:
                self.negative_words.add(row.strip())

    def count_neg_words(self, sentence):
        num_neg_words = 0
        for word in sentence.lower().split():
            if word in self.negative_words:
                num_neg_words += 1
        return num_neg_words
```

Load Features

```
lex_class = LexiconClassifier()

X_train_lexicon_features = []
X_test_lexicon_features = []

for string in X_text_test:
    X_test_lexicon_features.append([lex_class.count_neg_words(string)])

for string in X_text_train:
    X_train_lexicon_features.append([lex_class.count_neg_words(string)])
```

Combine RandomForestClassifier with LexiconClassifier

```
vec = CountVectorizer(ngram_range = (1,1))

X_train_w_lex = vec.fit_transform(X_text_train)
X_test_w_lex = vec.transform(X_text_test)

X_train_lexicon_features = np.array(X_train_lexicon_features)
X_test_lexicon_features = np.array(X_test_lexicon_features)

X_train_w_lex = sp.hstack((X_train_lexicon_features, X_train_w_lex))
X_test_w_lex = sp.hstack((X_test_lexicon_features, X_test_w_lex))
```



Model Selection Procedure

In our exploration of possible offensive language classification models, we determined:


- Random Forest + Lexicon machine learning model yielded the best validation results
- F1 Micro and F1 Macro scores of **0.8758** and **0.3113** respectively

Final Model Parameters + GridSearchCV Values


- `n_estimators = [10, 30, 50, 70, 90, 100]`
- `parameters = {'n_estimators': n_estimators}`
- `rand_forest_lex = RandomForestClassifier()`
- `clf_rand_lex = GridSearchCV(rand_forest_lex, parameters, cv = 10, verbose = 2)`
- `clf_rand_lex.fit(X_train_w_lex, y_train)`

Model Results - Classification Segmentation


--LinearSVC Prediction Results--
Offensive Language Segmentation



--RandomForest Prediction Results--
Offensive Language Segmentation



--RandomForest + Lexicon Prediction Results--
Offensive Language Segmentation





Model Results - Scoring

LinearSVC

F1 Micro: 0.7753
F1 Macro: 0.2911

Random Forest

F1 Micro: 0.8667
F1 Macro: 0.3095

Random Forest + Lexicon

F1 Micro: 0.8758
F1 Macro: 0.3113



Error Analysis - Random Forest + Lexicon

- **Step 1: Random sample of model results**
 - Our sample included 30 observations
- **Step 2: Manual review of 30 observations**
 - The model predicted 80% of the observations correctly (24 out of 30)
- **Step 3: Categorize Errors**
 - 6 observations were classified incorrectly
 - 5 were found to be *false negatives*
 - 1 was found to be a *false positive*

Rationale:

- Lack of specific negative words within the sample
- Model's inability to recognize the targeting of Twitter users in Tweets

Thank you!

